

# Mitigate Communication Barrier!

Recognise Sign Language using  
Neural Networks and Deep Learning

Presented by: Group 21

Jialin Liu (A0172938R)

Mansi Agarwal (A0218968J)

Nishtha Malhotra (A0228556W)

CS5242 Neural Networks and Deep Learning  
AY 2022/2023




**NUS**  
National University  
of Singapore

National University of Singapore

# Contents

- 1. Motivation, Problem Statement, Challenges**
- 2. Data Collection, Cleaning, Exploration**
- 3. Deep Learning Models Training and Evaluation**
  1. Methodology
  2. Image Preprocessing and Preparation
  3. Multiple Layer Perceptron (MLP)
  4. Convolutional Neural Network (CNN)
  5. Recurrent Neural Network (RNN)
- 4. Results Analysis & Interpretation**
- 5. Project Initiatives & Conclusion**
- 6. Future Developments**




# 1. Introduction

- ❖ Project Motivation
- ❖ Problem Statement
- ❖ Challenges



DREAM IT. BELIEVE  
IT. BUILD IT.



# Widespread Hearing Loss

Nearly **430 million** people suffer from deafness or hearing loss  
[Source: [WHO](#)]



Singapore Association for the Deaf (**SADeaf**) estimates **500,000** people suffer with hearing loss, only 1800 know sign language



THE SINGAPORE ASSOCIATION  
FOR THE DEAF  
SIGN LANGUAGE INTERPRETATION



# | Automatic Sign Language Recognition

## Significance


- Mitigate communication barrier for hard-of-hearing community with their First Responders and abled communities
- Create Educational resources, Work opportunities, Foster friendly professional environments
- Build easily accessible sign language instruction and specialized education programs



# Automatic Sign Language Challenges

## Challenges

- Multimodal: Signs are characterized by hand shapes, movements, facial expressions, body postures
- Similar Signs for different characters
- Variations due to orientation, hand shape of different people
- Characters with Dynamic Hand Gestures cannot be recognised by static image recognition systems




- ★ Similar signs with clenched fist for alphabets A, E, M, N, S, T
- ★ Similar signs for characters pairs '2' | 'V' and 'O' | '0'
- ★ Dynamic hand gestures characteristic to alphabets 'J' and 'Z'




# Problem Statement

**Build Supervised Learning  
American Sign Language (ASL)  
Recogniser System  
using deep neural networks  
capable of classifying static hand gesture  
images into corresponding 10 digit labels**



## 2. Data

- ❖ Collection
- ❖ Cleaning
- ❖ Exploration





# | Data Collection

## Image Scrapers

### Istockphoto and gettyimages


- BeautifulSoup: Obtain all image links using unique html element containing the image
- Different URLs can be provided to scraper to get images from URL link

### Google

- Used Google API [images-scraper](#) for collecting digit and character images from Google Images

Scrapped Images Statistics		
Istockphoto	GettyImages	Google
5000	5000	1440

## 3 Sources



# Data Cleaning and Exploration




## Cleaning


- Removed noisy images irrelevant to our use case of digit recognition

## Exploration

- Colour plays no role - Single Channel input images
- Challenges and Problems in collected images
  - Features difference in animated and human hand images
  - Imbalance sample size for 10 digits Label distribution
  - Needs standardisation of size and resolution



Distribution of sample size by digits as labels



# Training and Test Dataset


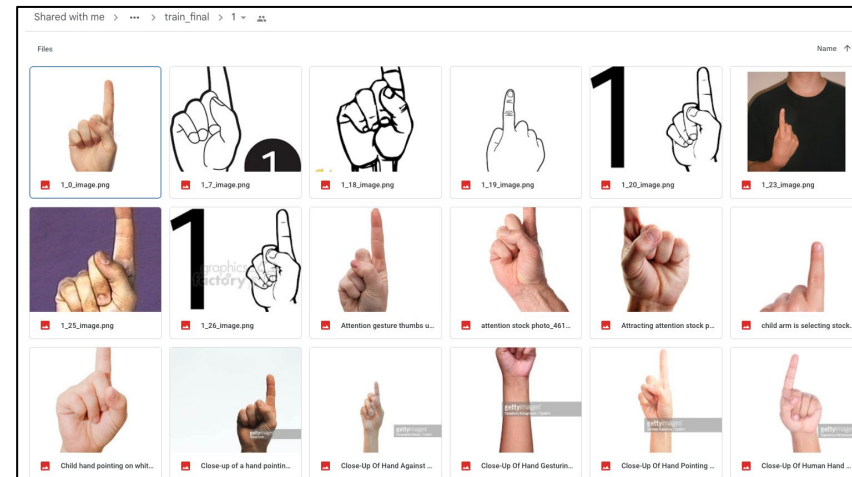
## Randomly split cleaned data into train and test

### Training dataset


- 672 Images
- Both human and animated hands

### Test dataset

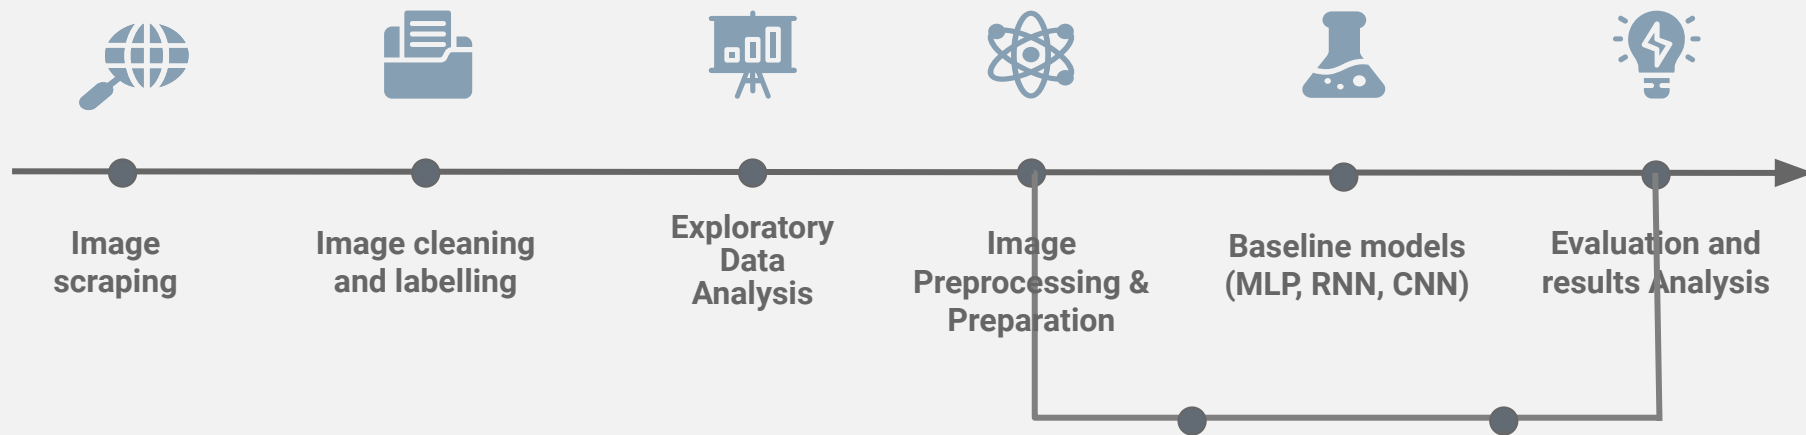
- 50 Images | 5 for each digit
- Only human hands more relevant to real-world use case



# 3. Deep Learning Models Training and Evaluation



# Methodology




## Standardised following considerations in all models


- **Setting random seed:** for reproducibility
- **Optimizer:** Adaptive learning rate algorithms like Adam.
- **Train batch size:** 16
- **Model hyperparameters:** learning rate, #epochs, #hidden layers, #hidden units

Data  
augmentation

Model Experiments  
for improvement








# Image Preprocessing and Preparation


## Image Transformation

- **Re-size:** 28x28, 56x56, 112x112, 224x224
- **Cropping from centre**
  - Given size as input, image cropped from centre
  - Automatically padded on edges for given size
- **Gray Scale:** Reduces to single channel
- **Feature Scaling and Normalisation:**
  - Converts PIL image to numpy array of range of [0,1]
  - Normalisation rescales images to have a mean of 0.485 and a standard deviation of 0.224

Original



Horizontal Flip



## Data Augmentation


### Advantages:

- Incorporates invariances and generalization
- Provides more samples for training


### Simulated real-world variances:

- Original Image transformations
- Horizontal Flip
- Original with Gaussian Noise
- Flipped with Gaussian Noise
- Original with Random Noise


Original with gaussian noise



Flipped with gaussian noise



Original with random erasing



# Multilayer Perceptron – *Model setup*


## Why MLP?

- Basic Deep Neural Network
- Consists series of fully connected layers

## Baseline model

- Input image normalisation
- 2 hidden layers with 120 and 84 neurons each
- Softmax output (10 classes)
- Relu activation function
- 40 epochs
- Adam optimizer

## Baseline MLP Architecture



```
MLP(  
  (h1): Linear(in_features=784, out_features=120, bias=True)  
  (h2): Linear(in_features=120, out_features=84, bias=True)  
  (bn1): BatchNorm1d(84, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (out): Linear(in_features=84, out_features=10, bias=True)  
)
```

# Multilayer Perceptron – *Experiments*

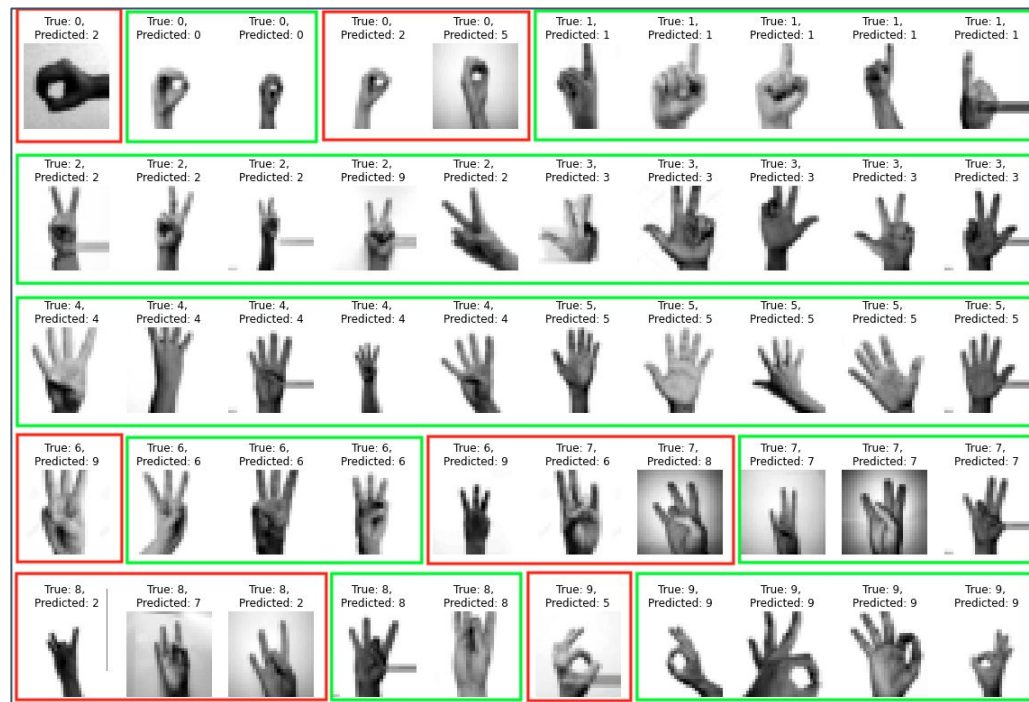

Metric	Experiment 1	Experiment 2	Experiment 3	Experiment 4	Experiment 5	Experiment 6	Experiment 7
Scenario	28x28 grayscale lr=0.001 Batch Size 16	28x28 grayscale With Data augmentation and Normalization Batch Size 16	56x56 grayscale With Data augmentation and normalisation Batch Size 16	28x28 grayscale With Data augmentation and Normalization Batch Size 16 higher/lower <b>learning rate</b>	28x28 grayscale With Data augmentation and Normalization Batch Size 16 more/less <b>neurons</b>	28x28 grayscale With Data augmentation and Normalization Batch Size 16 With <b>He</b> <b>initialisation</b>	28x28 grayscale With Data augmentation and <b>BatchNorm</b> and <b>batch size 32</b>
Validation Loss	1.75	0.7	0.3	1.9	0.9 / 0.5	0.5	0.2
Validation Accuracy	16%	52%	40%	10%	34% / 38%	48%	76%



# Multilayer Perceptron

## Best MLP Model Performance

- **Augmented Training dataset size:** 3360
- **Image size:** 28x28
- **Learning rate:** 0.001
- **Batch size:** 32
- **Num Epochs:** 40
- **Num Neurons:** [120, 84]
- **Batch Normalisation before output layer**



# **Multilayer Perceptron – *Limitations***

- Spatial information is lost when the image is flattened (matrix to vector) into an MLP
- Would not generalise well if the hand is at different part of the image as correlation of the image features (pixels) is not captured
- It includes too many parameters because it is fully connected. Hence more prone to overfitting
- The very dense web is formed can also resulting in redundancy and inefficiency.





# Convolutional Neural Network (CNN)

## Why CNN?

- Capable of interpreting spatial relationships between nearby pixels
- Accounts for image stationarity, applies same filter to different parts
- For detailed photos, CNN performs better than a MLP

## Baseline CNN Model Setup

- Original Image transformations
- Changed size to 112x112 for better resolution
- Train\_batch\_size = 16
- # epochs = 20, 50
- Learning\_rate = 0.0001

## Baseline CNN Model Performance

- Training starts to converge after 30 epochs
- Validation Accuracy range 78% - 84%

## Architecture

- ReLU activation function for Conv2D and FC1 layers
- Dropout Layers with  $p = 0.5$

```
CNN(  
  (conv1): Conv2d(1, 16, kernel_size=(5, 5), stride=(1, 1))  
  (conv2): Conv2d(16, 16, kernel_size=(5, 5), stride=(1, 1))  
  (conv3): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1))  
  (fc1): Linear(in_features=18432, out_features=256, bias=True)  
  (fc2): Linear(in_features=256, out_features=10, bias=True)  
)
```

### Baseline CNN Model Performance

Metric	Experiment 1 # epochs = 20	Experiment 2 # epochs = 50
Validation Loss	1.0989	1.3452
Validation Accuracy	70%	82%
Test Accuracy	70%	82%

# Convolutional Neural Network (CNN)

## CNN Augmented Data Model Setup

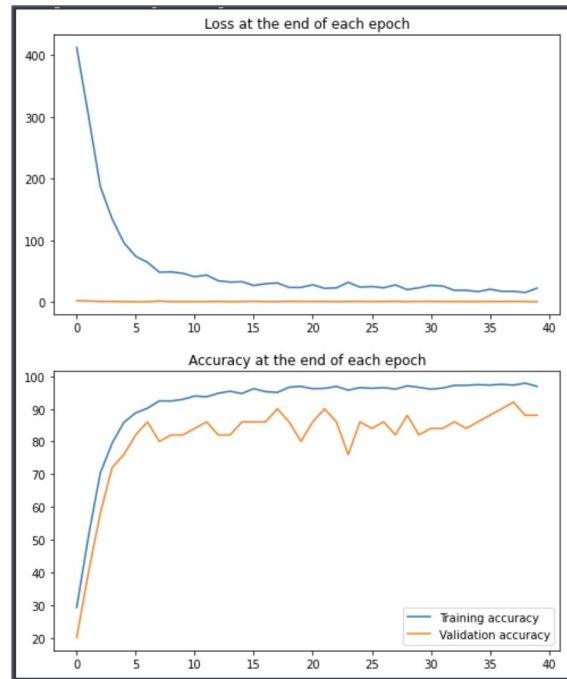
- 5 permutations to original image transformations
- Training data size = 3360 (672\*5 times)
- # epochs = 40 | learning\_rate = 0.0001 | batch\_size = 16
- Trained on same CNN Architecture

## CNN Model Performance with Augmented Data

- Training starts to converge after 7 epochs
- Validation Accuracy range 84% - 92%
- Performs better and converged faster than baseline CNN model

### CNN Model Performance with Augmented Data

Metric	Experiment 1 # epochs = 7	Experiment 2 # epochs = 50
Validation Loss	0.5109	0.5358
Validation Accuracy	86%	92%
Test Accuracy	78%	88%



# AlexNet CNN Architecture


## Why AlexNet Architecture?

- Deeper architecture with 8 layers: 5 Conv2D, 3 fully connected
- Overlapping pooling generally find it harder to overfit

## AlexNet Baseline Model Setup

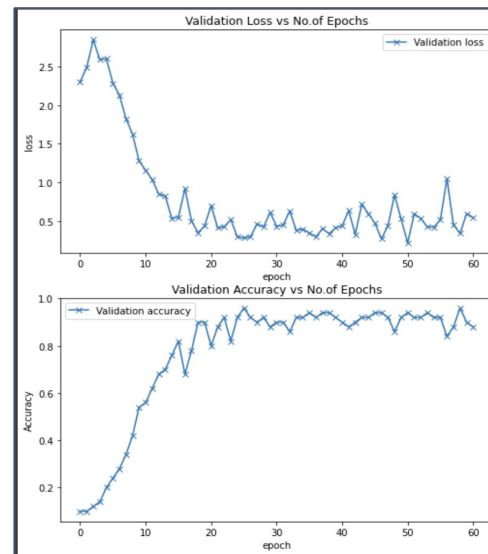
- Resize = 224 x 224 | Input channels = 3
- # epochs = 60 | learning\_rate = 0.0001 | batch\_size = 16

## AlexNet Architecture



## AlexNet Baseline Model Performance

- Training starts to converge at 30 epochs
- Validation Accuracy range 86% - 92%



# AlexNet CNN Architecture

## AlexNet Augmented Data Model Setup

- 5 transformations on original image transformations
- Training data size = 3360 (672\*5 times)
- # epochs = 40 | learning\_rate = 0.0001 | batch\_size = 16
- Addresses the problem of overfitting 60M parameters

## AlexNet Augmented Data Model Performance

- Training starts to converge within 10 epochs
- Validation Accuracy range 88% - 94%

Best CNN Model - AlexNet on Augmented Data		
Metric	Experiment 1 # training_data = 672	Experiment 2 # augmented_data = 3360
Validation Loss	0.3382	0.1173
Validation Accuracy	94% (34th epoch)	96% (epochs = 19th epoch)
<u>Test Accuracy</u>	<u>92%</u>	<u>94%</u>



# Recurrent Neural Network (RNN)


## Why RNN?

- Each image is treated as a sequence of 28 rows of 28 pixels (if the image size is 28). Hence RNN is used to study the sign language patterns in these sequences.

## Baseline model

- Input image normalisation
- 1 hidden layer RNN with 150 neurons
- 1 FC layer with output size of 10
- 30 epochs
- Adam optimizer


## Architecture





# Recurrent Neural Network (RNN)



Metric	Experiment 1	Experiment 2	Experiment 3	Experiment 4	Experiment 5	Experiment 6	Experiment 7
Scenario	28x28 grayscale lr=0.001	28x28 grayscale With Data augmentation Normalization Batch Size = 16	56x56 grayscale With Data augmentation Normalization	112x112 grayscale With Data augmentation Normalization	56x56 grayscale With Data augmentation With more neurons 250 Batch size = 32	28x28 grayscale With Data augmentation Batch Size = 32 Normalization Neurons = 150	28x28 grayscale With Data augmentation Batch Size = 32 Normalization Neurons = 250
Validation Loss	2.8	1.2	1.8	2.0	2.25	1.2	1.3
Validation Accuracy	21%	72%	42%	30%	22%	72%	66%




# Recurrent Neural Network (RNN)

## Image size 28x28



### Without Augmentation




### With Augmentation and Norm



### Doubled Batch Size





### Doubled Neurons




## Image size 56x56

### With Augmentation and Norm




### Doubled Batch Size & Neurons



# Recurrent Neural Network (RNN)

- Training dataset size (after data augmentation): 3360
- Image size: 28x28
- Lr: 0.001
- Batch size: 16
- Num Epochs: 30
- Num Neurons: 150
- Batch Normalisation before output layer



36 correct predictions




# Conclusion

**Deep Neural Network Models Performance Summary**

Model	AlexNet CNN	Multilayer Perceptron MLP	Recurrent NN RNN
Metric	Augmented Training dataset size: 3360 Image size: 3x224x224 Learning Rate: 0.0001 Batch Size: 16 Number of Epochs: 40 5 convolutional layers 3 fully connected layers Dropout, Adaptive Average Pooling	Augmented Training dataset size: 3360 Image size: 28x28 Learning Rate: 0.001 Batch size: 32 Num Epochs: 40 Num Neurons: [120, 84] Batch Normalisation before output layer.	Training dataset size (after data augmentation): 3360 Image size: 28x28 Learning Rate: 0.001 Batch size: 16 Num Epochs: 30 Num Neurons: [150] Batch Normalisation before output layer
Validation Loss	0.1173	0.2	1.2
Validation Accuracy	96% (epochs = 19)	76% (epochs = 40)	72% (epochs = 30)

# Project initiatives


- **3 image scrapers for 3 different data sources**
  - Can autogroup images into label folders
  - URL/Search based query
- **Analyse Epoch vs Validation Loss/Accuracy at every epoch**
- **Data Augmentation with 5 transformations**
- **Consolidated images to easily reflect correctly/wrongly labelled test data**





# Future Developments

- **Combine CNN-RNN**
- **Object Detection to detect the hand and signal**
- **Model to detect both the animated hands and human hands**



# Thank You



**NUS**  
National University  
of Singapore

National University of Singapore